Parallel scalable simulations of biological neural networks using TensorFlow: A beginner's guide

Saptarshi Soham Mohanta and Collins Assisi

Indian Institute of Science Education and Research, Pune, Maharashtra, India

*saptarshi.sohammohanta@students.iiserpune.ac.in
* collins@iiserpune.ac.in

Abstract

Neuronal networks are often modelled as systems of coupled, nonlinear, ordinary or partial differential equations. The number of differential equations used to model a network increases with the size of the network and the level of detail used to model individual neurons and synapses. As one scales up the size of the simulation it becomes important to use powerful computing platforms. Many tools exist that solve these equations numerically. However, these tools are often platform-specific. There is a high barrier of entry to developing flexible general purpose code that is platform independent and supports hardware acceleration on modern computing architectures such as GPUs/TPUs and Distributed Platforms. TensorFlow is a Python-based open-source package initially designed for machine learning algorithms, but it presents a scalable environment for a variety of computations including solving differential equations using iterative algorithms such as Runge-Kutta methods. In this article, organized as a series of tutorials, we demonstrate how to harness the power of TensorFlow's data-flow programming paradigm to solve differential equations. Our tutorial is a simple exposition of numerical methods to solve ordinary differential equations using Python and TensorFlow. It consists of a series of Python notebooks that accompany this paper. Over the course of five sessions, we lead novice programmers from writing programs to integrate simple 1-dimensional differential equations using Python, to solving a large system (1000's of differential equations) of coupled conductance-based neurons using a highly parallelised and scalable framework that uses Python and TensorFlow. Embedded in the tutorial is a physiologically realistic implementation of a network in the insect olfactory system. This system, consisting of multiple neuron and synapse types, can serve as a template to simulate other networks.

Motivation

Information processing in the nervous system spans a number of spatial and temporal scales. Millisecond fluctuations in ionic concentration at a synapse 10^- can cascade into long term (hours to days) changes in the behavior of the organism. Capturing the temporal scales and the details of the dynamics of the brain is a colossal computational endeavour. The dynamics of single and small networks of neurons can easily be simulated on a desktop computer with high level, readable, programing languages like Python. However, large networks of conductance based neurons are often simulated on clusters of CPUs. More recently, graphical processing units (GPUs) have become increasingly available due to lower (though still prohibitive) costs and from cloud services like Google Cloud and Amazon AWS among others. Writing code for each of

Python (NumPy) vs TensorFlow Benchmarks

A. Matrix Multiplication   *B. Numerical Integration*

Fig 1. Comparison of Python and Tensorflow on an 8 core, 2.1GHz desktop with 2 Intel Xeon E5-2620 processors

these platforms is non trivial and requires a considerable investment of time to master different software tools. For example, implementing parallelism in multiple core shared memory architectures is often achieved using Open Multi-Processing (OpenMP) with C or C++. Message Passing Interface (MPI) libraries are used to implement code that computes over high performance computing clusters. Compute Unified Device Architecture (CUDA) allows users to run programs on NVIDIA's GPUs. However, code written for one platform cannot be used on other platforms. This poses a high barrier of entry for computational neuroscientists conversant with a high-level programing language, attempting to test simulations on different platforms.

This is where TensorFlow, a Python library that was originally written to cater to machine learning applications [], comes in handy. TensorFlow is highly scalable. Code written using TensorFlow functions can work seamlessly on single cores, multi-core shared memory processors, high-performance computer clusters, GPUs and Tensor processing units (TPUs - a propreitary chip designed by Google to work with TensorFlow). We found (as others have (ref??)), that TensorFlow functions can be used to implement numerical methods to solve ODEs. Doing so gave us a significant speed-up even on a single desktop with a multicore processor compared to similar Python code that did not use TensorFlow functions and operated on a single core. The code itself was highly readable and could be debugged with ease. Familiarity with the Python programming language and a brief introduction to some TensorFlow functions proved sufficient to write the code. Python is a an extremely popular programming language that used across number of disciplines and has found a broad user base among Biologists. We found that introducing a few TensorFlow constructs to Python, an easy addition to a familiar language, can bring readers to a point where they can simulate large networks of neurons in a platform independent manner. Further, by piggybacking on TensorFlow we will also be able to take advantage of an active TensorFlow developer community in addition to a wide range of Python libraries that are already available.

Note on GPU/TPU cloud

These tutorials were written to address a group of undergraduate students in our institute. These students came from diverse backgrounds and had a basic introduction to Python during their first semester. They were interested in working on problems in Computational Neuroscience. Our goal was to introduce them to some of the numerical tools and mathematical models in Neuroscience while also allowing them to tinker with advanced projects that required writing code that ran on distributed systems. We were careful to keep the innards of the code visible —the form of the integrator, the specification of the differential equations and the ability to modify the code to suit their needs was particularly important. Towards the end of the tutorial, that many students managed to devour within a day or two, they were in a position to write codes simulating networks of neurons in the antennal lobe, firing rate models of grid cells, detailed networks of stellate cells and inhibitory interneurons and synaptic plasticity.

How to use this Tutorial

A reader familiar with Python will find this tutorial accessible. We use a number of Numpy and Matplotlib functions to simulate and display figures. These libraries are well documented with excellent introductory guides. During Day 1 of this tutorial we introduce numerical integration using Python without using any TensorFlow functions. On Day 2 we use TensorFlow functions to implement the integrator. Day 5 talks about memory management in TensorFlow. On Days 3 and 4 we use the code developed on Days 1 and 2 to simulate networks of conductance based neurons. eaders who are not interested in simulating neurons, but solving differential equations in other domains, Days 1, 2 and 5 are self-contained and can be read without reference to the material covered on Days 3 and 4. Our simulations were run on a Linux desktop running Ubuntu xx.xx and on Google Cloud. We recommend that readers install Python 3.6 or above, Jupyter Notebook, Numpy Python package, Matplotlib Python package, and TensorFlow 1.13 or above using the Anaconda distribution of Python 3. The tutorials are linked in the supplementary material as Python notebooks (.ipynb files) that can be accessed using Jupyter and as .html files that can be read using any browser.

Day 1: Solving ODEs using Python

In this tutorial we are interested in solving ordinary differential equations of the form,

$$\frac{dx}{dt} = f(x, t) \tag{1}$$

where, x is an $N-$dimensional vector and t typically stands for time. The function $f(x, t)$ may be a nonlinear function of x that explicitly depends on t. In addition to specifying the form of the differential equation, one must also specify the value of x at a particular time point. Say, $x = x_0$ at $t = t_0$. It is often not possible to derive a closed form solution for 1. Therefore numerical methods to solve these equations are of great importance. One such example at the core of our tutorial are the Hodgkin-Huxley equations (ref) describing the dynamics of action potential generation and propagation in the giant axon of the squid. Alan Hodgkin and Andrew Huxley arrived at these equations after a series of clever experiments that tested the limits of experimental technology available at the time. It also tested the limits of computational tools available. The form of the differential equations they derived contained nonlinearities that made it analytically intractable. In order to compute action potentials, Huxley numerically intergrated the equations using a hand operated Bunsviga mechanical calculator. The calculation took nearly three weeks to complete (ref). They used a

numerical method due to Hartree (ref 1932)integrate the differential equations. Each
iteration that calculated the value of the solution at subsequent time points consisted of
9 steps. In calculating the solution over time, they also varied the step size such that
dynamics occurring over faster time scales (such as the rising phase of the action
potential) were calculated with time step sizes of $0.1 - 0.2$ms while slower dynamics
(such as the small highly damped oscillation following a spike) was calculated with time
steps that were an order of magnitude higher (1ms).

   In this tutorial, we illustrate two numerical methods to iteratively compute each
time step of the solution. The first is a simple one-step method known as the Euler's
method of integration. The second is another popular method, the Runge-Kutta
method of order 4 (abberviated as RK4) that is more accurrate than the Euler's
method, but requires additional computations to calculate subsequent time steps.

Euler's Method for Numerical Integration

Our goal is to solve 1 - calculate $x(t)$ given an initial condition $x(t_n) = x_n$. The
simplest method to solve 1 numerically is the Euler's method. Here we start from
$x(t = t_0) = x_0$ and compute the solution at subsequent time points
$(t_0 + \epsilon, t_0 + 2\epsilon, t_0 + 3\epsilon \dots)$ iteratively. Each step of the computation is done using the
same formula that can be derived by truncating a Taylor series after the first term.
That is, the solution at time $t + \epsilon$ is given by,

$$x(t + \epsilon) = x_0 + \epsilon \frac{dx}{dt} + \mathcal{O}(\epsilon^2) \tag{2}$$

where $\frac{dx}{dt} = f(x, t)$. The higher order terms $\mathcal{O}(\epsilon^2)$ are ignored in this approximation.


Fig 2. Euler's method

   A geometric interpretation of the Euler's method is shown in figure 2. The solution
of a particular differential equation is represented by the solid blue line in the figure.
The dashed line approximates this solution by iteratively calculating x at different time
points using equation 2. If the value of the solution at t_n is x_n, $f(x_n, t)$ is the slope of

the tangent to the solution at t_n. If ϵ is sufficiently small, the solution at $t_{n+1} = t_n + \epsilon$ can be approximated by linearly extrapolating from x_n to $x_n + \epsilon f(x_n, t_n)$

Implementation of Euler's Method in Python

Let $\frac{dx}{dt} = 5x$. We wish to calculate $x(t)$ over the interval $t \in [0, 2)$ given the initial condition $x(0) = 1$. The exact solution of this equation is $x(t) = e^{5t}$. In our implementation of Euler's method, we used the Python library Numpy to create and operate on arrays, and the plotting library Matplotlib, to display the results. We implement the Euler's Method in Python as follows:

```python
import numpy as np
import matplotlib.pyplot as plt
def f(x,t): # define the function f(x,t)
    return 5*x
epsilon = 0.01 # define timestep
t = np.arange(0,2,epsilon) # define an array for t
x = np.zeros(t.shape) # define an array for x
x[0]= 1 # set initial condition
for i in range(1,t.shape[0]):
    x[i] = epsilon*f(x[i-1],t[i-1])+x[i-1] # Euler Integration Step
```


Fig 3. Comparison of actual solution and approximation by Euler's method

The exact solution and the numerical solution are compared in Fig 3. Notice that the approximation begins to diverge from the actual solution of the equation as the number of iterations increase. The omission of terms $\mathcal{O}(\epsilon^2)$ leads to a truncation error per step that can accumulate over time.

Implementing Euler's method to solve a system of differential equations

Euler's Method can also be easily applied to systems of equations. The Initial Value problem now becomes:

$$\frac{d\vec{X}}{dt} = \vec{f}(\vec{X}, t) \tag{3}$$

$$\vec{X}(t_o) = \vec{X_o} \tag{4}$$

where $\vec{X} = [X_1, X_2...]$ and $\vec{f}(\vec{X}, t) = [f_1(\vec{X}, t), f_2(\vec{X}, t)...]$. We rewrite Euler's
method as:

$$t_{n+1} = t_n + \epsilon \tag{5}$$

$$\vec{X}(t_{n+1}) = \vec{X}(t_n) + \epsilon\vec{f}(\vec{X}(t_n), t_n) \tag{6}$$

Let $\frac{d\vec{X}}{dt} = f(\vec{X}, t)$, we wish to find $\vec{X}(t)$ over $t \in [0, 2)$, given that $\vec{X}(t) = [x, y]$,
$\vec{X}(0) = [1, 0]$ and $f(\vec{X}, t) = [x - y, y - x]$. We implement the modified algorithm as
follows:

```
1  def f(X,t): # the function f(X,t) now takes a vector X as input
2      x,y = X #the first and the second elements of X are assigned to x and y
3      return np.array([x-y,y-x])
4  t = np.arange(0,2,epsilon) # define an array for t
5  X = np.zeros((2,t.shape[0])) # initialize an array for X
6  X[:,0]= [1,0] # set initial condition
7  for i in range(1,t.shape[0]):
8      X[:,i] = epsilon*f(X[:,i-1],t[i-1])+X[:,i-1] # Euler Integration Step
```

A generalized code implementing the Euler method

Here we rewrite the code in a modular fashion and cast the integrator as a function that
takes in 3 inputs ie. the function $\vec{f}(\vec{y}, t)$ where $\frac{d\vec{y}}{dt} = f(\vec{y}, t)$, the time array, and initial
vector $\vec{y}_0$. We will find this form to be particularly useful when we use tensorflow
constructs to code the integrator. Further, it allows us to write multiple integrating
functions (for example Euler or RK4) within the same class and call a specific integrator
as needed. In addition we also introduce a function to ensure that the correct inputs are
given to the integrator failing which an error message is generated.

Algorithm

- Get the required inputs: function $\vec{f}(\vec{y}, t)$, initial condition vector $\vec{y}_0$ and time
  series t. Entering a time series t allows for greater control over ϵ as it can now
  vary for each timestep.

- Check if the input is of the correct datatype ie. floating point decimal.

- Create a zero matrix to hold the output.

- For each time step, perform the update $\vec{y}$ using the Euler method with variable ϵ
  and store it in the output matrix.

- Return the output time series [number of equations $\times$ iterations] matrix.

```
1   def check_type(y,t): # Ensure Input is Correct
2       return y.dtype == np.floating and t.dtype == np.floating
3   class _Integrator():
4       def integrate(self,func,y0,t):
5           time_delta_grid = t[1:] - t[:-1]
6           y = np.zeros((y0.shape[0],t.shape[0]))
7           y[:,0] = y0
8           for i in range(time_delta_grid.shape[0]):
9               y[:,i+1]= time_delta_grid[i]*func(y[:,i],t[i])+y[:,i]
10          return y
11  def odeint_euler(func,y0,t):
12      y0 = np.array(y0)
13      t = np.array(t)
14      if check_type(y0,t):
15          return _Integrator().integrate(func,y0,t)
16      else:
17          print("error encountered")
18  solution = odeint_euler(f,[1.,0.],t)
```

Runge-Kutta Methods for Numerical Integration

Euler's method $x_{n+1} = x_n + \epsilon f(x_n, t_n)$ calculates the solution at $t_{n+1} = t_n + \epsilon$ given the solution at t_n. In doing so we use the derivative at t_n though its value may change throughout the interval $[t, t + \epsilon]$. This results in an error in the order of $\mathcal{O}(\epsilon^2)$. By calculating the derivatives at intermediate steps, one can reduce the error at each step. Consider the following second order method where the slope is calculated at t_n and $t_n + \frac{\epsilon}{2}$.

$$k_1 = \epsilon f(x_n, t_n) \tag{7}$$

$$k_2 = \epsilon f(x_n + \frac{k_1}{2}, t_n + \frac{\epsilon}{2}) \tag{8}$$

$$x_{n+1} = x_n + k_2 + O(\epsilon^3) \tag{9}$$

This method is called the second order Runge-Kutta method or the midpoint method. Figure 4 is a schematic description of the second order Runge-Kutta methd. The blue curve denotes a solution of some differential equation. The goal is to calculate

But we do not have to stop here. By further rewriting the equation, we can cancel higher order error terms and reach the most commonly used fourth-order Runge-Kutta Methods or RK4 method, which is described below:

$$k_1 = f(x_n, t_n) \tag{10}$$

$$k_2 = f(x_n + \epsilon \frac{k_1}{2}, t_n + \frac{\epsilon}{2}) \tag{11}$$

$$k_3 = f(x_n + \epsilon \frac{k_2}{2}, t_n + \frac{\epsilon}{2}) \tag{12}$$

Fig 4. Second order Runge-Kutta method

$$k_4 = f(x_n + \epsilon k_3, t_n + \epsilon) \tag{13}$$

$$y_{n+1} = y_n + \frac{\epsilon}{6}(k_1 + 2k_2 + 2k_3 + k_4) + \mathcal{O}(\epsilon^5) \tag{14}$$

Note that this numerical method is again easily converted to a vector algorithm by simply replacing x_i by the vector $\vec{X}_i$. This method is what we will use to simulate our networks.

Generalized RK4 Method in Python

We can now modify the Euler Integration code implemented earlier iwith a generalized function for RK4 that takes three inputs —the function $f(\vec{y}, t)$ when $\frac{d\vec{y}}{dt} = f(\vec{y}, t)$, the time array, and an initial vector $\vec{y_0}$. The code can be updated as follows,

```
1  # RK4 Integration Steps replace the Euler's Updation Steps
2  k1 = func(y[:,i], t[i])
3  half_step = t[i] + time_delta_grid[i] / 2
4  k2 = func(y[:,i] + time_delta_grid[i] * k1 / 2, half_step)
5  k3 = func(y[:,i] + time_delta_grid[i] * k2 / 2, half_step)
6  k4 = func(y[:,i] + time_delta_grid[i] * k3, t + time_delta_grid[i])
7  y[:,i+1]= (k1 + 2 * k2 + 2 * k3 + k4) * (time_delta_grid[i] / 6) + y[:,i]
```

As an **Exercise**, try to solve the equation of a simple pendulum and observe its dynamics using Euler Method and RK4 methods. The equation of motion of a simple pendulum is given by:

$$\frac{d^2 s}{dt^2} = L\frac{d^2\theta}{dt^2} = -g\sin\theta \tag{15}$$

where L = Length of String and θ = angle made with vertical. To solve this second order differential equation you may use a dummy variable ω representing angular velocity such that:

$$\frac{d\theta}{dt} = \omega \tag{16}$$

$$\frac{d\omega}{dt} = -\frac{g}{L}\sin\theta \tag{17}$$

Day 2: Let the Tensors Flow!

An Introduction to TensorFlow

TensorFlow is an open-source library that was developed by researchers and engineers in the Google Brain team. TensorFlow has a number of functions that make it particularly suitable for machine learning applications. However, TensorFlow is primarily an interface for numerical computation (ref:TensorFlow whitepaper). All computations in TensorFlow are specified as directed graphs (nodes connected by arrows) known as data flow graphs. Nodes are operations such as addition, multiplication etc. The incoming edges for each node are tensors (scalars, vectors, matrices and higher dimensional arrays) that are the actual values that are operated upon. The output is also a tensor that results from computation. For example, consider the following computation where two vectors a and b serve as inputs to the node, a matrix multiplication operation, that produces a matrix c as output.


Fig 5. Example of a simple computational graph

The following program implements the computation described in 5.

```
# Creating nodes in the computation graph
a = tf.constant([[1.],[2.],[3.]], dtype=tf.float64) # a 3x1 column matrix
b = tf.constant([[1.,2.,3.]], dtype=tf.float64) # a 1x3 row matrix
```

```
4   c = tf.matmul(a, b)
5   # To run the graph, we need to create a session.
6   # Creating the session initializes the computational device.
7   sess = tf.Session() # start a session
8   output = sess.run(c) # compute the value of c
9   sess.close() # end the session
10  print(output)
11  # TO automatically close the session after computation, Use:
12  # with tf.Session() as sess:
13  #     output = sess.run(c)
```

By specifying the computation graph we also specify the dependcies among the nodes. 195
One can thus split the graph into smaller chunks or sub-graphs that can be 196
independently computed by different devices that coordinate with each other. This 197
makes it possible to develop programs that are device independent and scalable across 198
CPUs, GPUs , TPUs and clusters of servers. 199

Efficient recursion in TensorFlow 200

Numerical integration is essentially a recursive process over a time array $[t_0, t_1, t_2, \ldots t_n]$. 201
The updation rules for both the Euler and the RK4 integrators can just be written as a 202
recursive function F such that $X_{i+1} = F(X_i, t_i, \epsilon_i)$. The solution of the differential 203
equation is the array $[X_0, F(X_0, t_0, \epsilon_0), F(F(X_0, t_0, \epsilon_0), t_1, \epsilon_1)...]$. We us the TensorFlow 204
function `tf.scan` to iterate over the time array. The arguments of `tf.scan` are a 205
recursive function, the list to iterate over and the initial value. If the initial value is not 206
specified `tf.scan` uses the first element of the list as an initializer. As an example, 207
consider the following program that calculates the cumulative sum over a list. Every 208
step involves adding an element from the list onto the last addition. 209

```
1   # define the recursive function that takes in two values the
2   # accumulated value and the additional input from a list.
3   def recursive_addition(accumulator,new_element):
4       return accumulator+new_element
5   # define the list over which we iterate
6   elems = np.array([1, 2, 3, 4, 5, 6])
7   # tf.scan takes in three inputs: the recursive function, the
8   # list to iterate over and the initial value. If an initial
9   # value is not provided, its taken as the first element of elems.
10  # accumulate with no initializer
11  cum_sum_a = tf.scan(recursive_addition, elems)
12  # accumulate with initializer as the number 5
13  cum_sum_b = tf.scan(recursive_addition, elems, tf.constant(5,dtype=tf.int64))
14  with tf.Session() as sess:
15      output_a = sess.run(cum_sum_a)
16      output_b = sess.run(cum_sum_b)
17  print(output_a)
18  print(output_b)
19  # This prints :
20  #[ 1   3   6 10 15 21]
21  #[ 6   8 11 15 20 26]
```

Exercise Use `tf.scan()` to compute the Fibonacci sequence. 210

Euler Integration Function in TensorFlow

We now implement Euler's method using `tf.scan` to iterate over the time array. Note
that the function `scan_func` that defines each step of Euler's method, is now an input
to `tf.scan`.

```python
def tf_check_type(t, y0): # Ensure Input is Correct
    if not (y0.dtype.is_floating and t.dtype.is_floating):
        # The datatype of any tensor t is accessed by t.dtype
        raise TypeError('Error in Datatype')
class _Tf_Integrator():
    def integrate(self, func, y0, t):
        time_delta_grid = t[1:] - t[:-1]
        def scan_func(y, t_dt):
            t, dt = t_dt
            dy = dt*func(y,t)
            return y + dy
        # iterating over (a,b) where a and b are lists of same size
        # results in the ith accumulative step in tf.scan receiving
        # the ith elements of a and b zipped together
        y = tf.scan(scan_func, (t[:-1], time_delta_grid),y0)
        return tf.concat([[y0], y], axis=0)
def tf_odeint_euler(func, y0, t):
    # Convert input to TensorFlow Objects
    t = tf.convert_to_tensor(t, preferred_dtype=tf.float64, name='t')
    y0 = tf.convert_to_tensor(y0, name='y0')
    tf_check_type(y0,t)
    return _Tf_Integrator().integrate(func,y0,t)

# Define a function using Tensorflow math operations.
# This creates the computation graph.
def f(X,t):
    # extracting a single value eg. X[0] returns a single value but
    # we require a tensor, so we extract a range with one element.
    x = X[0:1]
    y = X[1:2]
    out = tf.concat([x-y,y-x],0)
    return out
y0 = tf.constant([1,0], dtype=tf.float64)
epsilon = 0.01
t = np.arange(0,2,epsilon)
# Define the final value (output of scan) that we wish to compute
state = tf_odeint_euler(f,y0,t)
# Start a TF session and evaluate state
with tf.Session() as sess:
    state = sess.run(state)
```

RK4 Integration Function in TensorFlow

Now, we implement the RK4 integrator. Note that here we replace the single step
iterator used for the Euler's with a four step RK4 iterator. In addition, to make the
code more modular, we define a function `_step_func()` that is called by `scan_func` and
calculates the next step of the RK4 integrator. The rest of the program remains the
same as the Euler's method implemented above.

```
1   def integrate(self, func, y0, t):
2           time_delta_grid = t[1:] - t[:-1]
3           def scan_func(y, t_dt):
4               t, dt = t_dt
5               dy = self._step_func(func,t,dt,y) # Make code more modular.
6               return y + dy
7           y = tf.scan(scan_func, (t[:-1], time_delta_grid),y0)
8           return tf.concat([[y0], y], axis=0)
9       def _step_func(self, func, t, dt, y):
10          k1 = func(y, t)
11          half_step = t + dt / 2
12          dt_cast = tf.cast(dt, y.dtype) # Failsafe
13          k2 = func(y + dt_cast * k1 / 2, half_step)
14          k3 = func(y + dt_cast * k2 / 2, half_step)
15          k4 = func(y + dt_cast * k3, t + dt)
16          return tf.add_n([k1, 2 * k2, 2 * k3, k4]) * (dt_cast / 6)
```

Exercise, Simulate the non-linear Lorentz Attractor using Euler Method and RK4 on TensorFlow which is given by the equations:

$$\frac{dx}{dt} = \sigma(y - x) \tag{18}$$

$$\frac{dy}{dt} = x(\rho - z) - y \tag{19}$$

$$\frac{dz}{dt} = xy - \beta z \tag{20}$$

Use the values $\sigma = 10$, $\beta = \frac{8}{3}$, $\rho = 28$. Simulate these equations for similar initial conditions and compare how the trajectories diverge. The solution of the Lorenz equations should resemble Figure 6

Lorenz Attractor Simulation

Fig 6. Example of a simple computational graph

Day 3: Neurons in Silicon

The electric potential measured across the membranes of excitable cells, such as neurons or heart cells, can undergo transient changes when perturbed by external inputs. When the inputs to a neuron are sufficiently large these transient changes can regeneratively build up into a large deviation from the resting state known as an action potential. Action potentials propagate along the axon, largely undiminished, and perturb post-synaptic neurons. The Hodgkin-Huxley model is a system of differential equations that describe the generation an action potential and its propagation along the axon.

What is the Hodgkin Huxley Neuron Model?

The cell membrane, a 5nm thick lipid bilayer, separates the inside from the outside of the neuron. The membrane is largely impermeable to charged ions present on either side. The concentration of Na^+ ions outside the cell is greater than its concentration inside, while K^+ ions are relatively abundant inside compared to the outside. In addition to these there are chloride (Cl^-), calcium (Ca^{2+}) and magnesium ions (Mg^+) that populate the cellular milieu. The differences in ionic abundances across the membrane causes a net accumulation of positive ions on one side of the membrane and negative ions on the other, and thus a potential difference across the membrane. Embedded on the membrane are ion channels that are highly selective to the ion species it lets across. In the squid axon, Hodgkin and Huxley found that there were only two types of ion channels (Na^+ and K^+), in addition to a non-specific leak channel that, most likely, allowed the passage of Cl^- ions. The Hodgkin-Huxley model of neurons can be understood with the help of an equivalent electrical circuit 7) The cell membrane acts as a capacitor. An total injected current (I) can be written as the sum of the

capacitive current I_C, ionic currents I_{Na} and I_K and the leak current I_L.
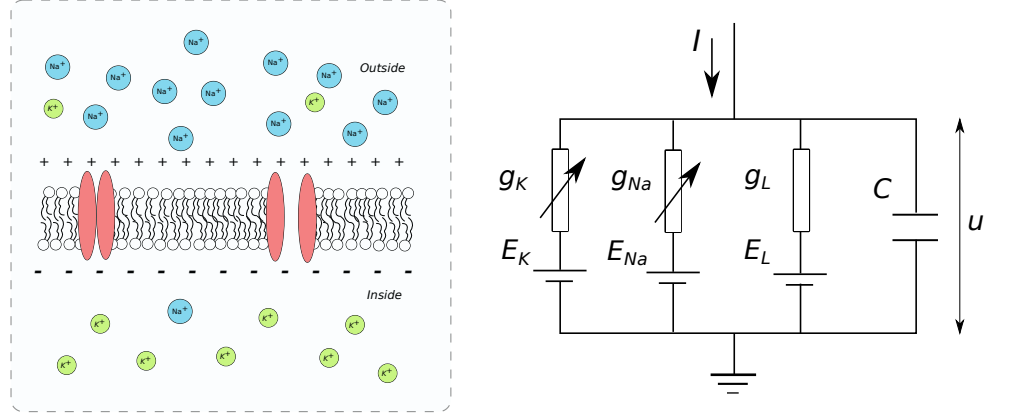
Fig 7. HH Neuron

$$I = I_C(t) + I_{Na}(t) + I_K(t) \tag{21}$$

where,

$$C_m = 1\mu F/cm^2 \tag{22}$$
$$I_{Na} = g_{Na}(u - E_{Na}) \tag{23}$$
$$I_K = g_k(u - E_K)I_L = g_L(u - E_L) \tag{24}$$

The equation describing the membrane potential can thus be written as follows,

$$C_m\frac{du}{dt} = -I_{Na}(t) - I_K(t) - I_L(t) + I(t) \tag{25}$$

Hodgkin and Huxley discovered that the Na and the K channels do not act as Ohmic conductances, but are modulated by the potential across the membrane. Changes in potential had a nonlinear effect on flow of ionic currents. They fitted their experimental results to a system of differential equations that described the temporal evolution of the membrane potential in terms of changes in ionic currents (chiefly Na$^+$ and $^+$).

Hogkin and Huxley found the Na and K ion currents to be dependent on the voltage and of the form given below:

$$I_{Na} = g_{Na}m^3h(u - E_{Na}) \tag{26}$$

$$I_K = g_K n^4(u - E_K) \tag{27}$$

$$I_L = g_L(u - E_L) \tag{28}$$

where $E_{Na} = 50\ mV$, $E_K = -95\ mV$ and $E_L = -55\ mV$ are the reversal potentials; $g_{Na} = 100\ \mu S/cm^2$, $g_K = 10\ \mu S/cm^2$ and $g_L = 0.15\ \mu S/cm^2$ are the channel conductances; and m,h, and n are gating variables that follow the dynamics given by:

$$\frac{dm}{dt} = -\frac{1}{\tau_m}(m - m_0) \tag{29}$$

$$\frac{dh}{dt} = -\frac{1}{\tau_h}(h - h_0) \tag{30}$$

$$\frac{dn}{dt} = -\frac{1}{\tau_n}(n - n_0) \tag{31}$$

where τ_m, τ_h and τ_n are voltage dependent time constants and m_0, h_0 and n_0 are voltage dependent asymptotic gating values. These functions are empirically determined for different types of neurons. For an example, take a look at figure.

Implementing the Dynamical Function for an Hodkin Huxley Neuron

For integration, we will use the TensorFlow RK4 integrator that we created. A simple Hodgkin Huxley Neuron has a 4 main dynamical variables: V = Membrane Potential, m = Sodium Activation Gating Variable, h = Sodium Inactivation Gating Variable, and n = Potassium Channel Gating Variable. And the dynamics are given by Eq (**??**), Eq (29),Eq (30) and Eq (31) where the values of τ_m, τ_h, τ_n, m_0, h_0, n_0 are given empirically derived formulae (Fig) and channel currents are determined by Eq (26),Eq (27) and Eq (28).

```
# Step 1: Defining Parameters of the Neuron
C_m = 1
g_K = 10
E_K = -95
g_Na = 100
E_Na = 50
g_L = 0.15
E_L = -55
# Step 2: Defining functions to calculate tau_x and x_0
# Note: Always use TensorFlow functions for all operations.
def K_prop(V):
    T = 22
    phi = 3.0**((T-36.0)/10)
    V_ = V-(-50)
    alpha_n = 0.02*(15.0 - V_)/(tf.exp((15.0 - V_)/5.0) - 1.0)
    beta_n = 0.5*tf.exp((10.0 - V_)/40.0)
    t_n = 1.0/((alpha_n+beta_n)*phi)
    n_0 = alpha_n/(alpha_n+beta_n)
    return n_0, t_n
def Na_prop(V):
    T = 22
    phi = 3.0**((T-36)/10)
    V_ = V-(-50)
    alpha_m = 0.32*(13.0 - V_)/(tf.exp((13.0 - V_)/4.0) - 1.0)
    beta_m = 0.28*(V_ - 40.0)/(tf.exp((V_ - 40.0)/5.0) - 1.0)
    alpha_h = 0.128*tf.exp((17.0 - V_)/18.0)
    beta_h = 4.0/(tf.exp((40.0 - V_)/5.0) + 1.0)
```

```
28      t_m = 1.0/((alpha_m+beta_m)*phi)
29      t_h = 1.0/((alpha_h+beta_h)*phi)
30      m_0 = alpha_m/(alpha_m+beta_m)
31      h_0 = alpha_h/(alpha_h+beta_h)
32      return m_0, t_m, h_0, t_h
33  # Step 3: Defining function that calculate Neuronal currents
34  def I_K(V, n):
35      return g_K  * n**4 * (V - E_K)
36  def I_Na(V, m, h):
37      return g_Na * m**3 * h * (V - E_Na)
38  def I_L(V):
39      return g_L * (V - E_L)
40  # Step 4: Define the function dX/dt where X is the State Vector
41  def dXdt(X, t):
42      V = X[0:1]
43      m = X[1:2]
44      h = X[2:3]
45      n = X[3:4]
46      dVdt = (5 - I_Na(V, m, h) - I_K(V, n) - I_L(V)) / C_m
47      # Here the current injection I_injected = 5 uA
48      m0,tm,h0,th = Na_prop(V)
49      n0,tn = K_prop(V)
50      dmdt = - (1.0/tm)*(m-m0)
51      dhdt = - (1.0/th)*(h-h0)
52      dndt = - (1.0/tn)*(n-n0)
53      out = tf.concat([dVdt,dmdt,dhdt,dndt],0)
54      return out
55  # Step 5: Define Initial Condition and Integrate
56  y0 = tf.constant([-71,0,0,0], dtype=tf.float64)
57  epsilon = 0.01
58  t = np.arange(0,200,epsilon)
59  state = odeint(dXdt,y0,t)
60  with tf.Session() as sess:
61      state = sess.run(state)
```

Simulating Multiple Independent HH Neurons at the Same Time

Although, simulating a Single Hodgkin-Huxley Neuron is possible in TensorFlow, the real ability of tensorflow can be seen only when a large number of simultaneous diffential equations are to be solved at the the same time. Let's try to simulate 20 independent HH neurons with different input currents and characterise the firing rates.

Methods of Parallelization

TensorFlow has the intrinsic ability to speed up any and all Tensor computations using available multi-cores, and GPU/TPU setups. There are two major parts of the code where TensorFlow can help us really speed up the computation:

1. **RK4 Steps:** Since the TensorFlow implementation of the Integrator utilizes Tensor calculations, TensorFlow will automatically speed it up.

2. **Functional Evaluations:** Looking at Dynamical Equations that describe the neuronal dynamics, its easy to notice that all simple HH Neurons share the same

or atleast similar dynamical equations but will vary only in the values of parameters. We can exploit this to speed up the computations.

Say $\vec{X} = [V, m, n, h]$ is the state vector of a single neuron and its dynamics are defined using parameters $C_m, g_K, ...E_L$ equations of the form:

$$\frac{d\vec{X}}{dt} = [f_1(\vec{X}, C_m, g_K, ...E_L), f_2(\vec{X}, C_m, g_K, ...E_L)...f_m(\vec{X}, C_m, g_K, ...E_L)] \qquad (32)$$

We have to somehow convert these to a form in which all evaluations are done as vector calculations and NOT scalar calculations.

So, what we need for a system of n neurons is to have a method to evaluate the updation of $\mathbf{X} = [\vec{X_1}, \vec{X_2}...\vec{X_n}]$ where $\vec{X_i} = [V_1, m_1, n_1, h_1]$ is the state vector of the ith neuron. Now there is a simple trick that allows us to maximize the parallel processing. Each neuron represented by $\vec{X_i}$ has a distinct set of parameters and differential equations.

Now, despite the parameters being different, the functional forms of the updation is similar for the same state variable for different neurons. Thus, the trick is to reorganize $\mathbf{X}$ as $\mathbf{X'} = [(V_1, V_2, ...V_n), (m_1, m_2, ...m_n), (h_1, h_2, ...h_n), (n_1, n_2, ...n_n)] = [\vec{V}, \vec{m}, \vec{h}, \vec{n}]$. And the parameters as $\vec{C_m}, \vec{g_K}$ and so on.

Now that we know the trick, what is the benefit? Earlier, each state variable (say V_i) had a DE of the form:

$$\frac{dV_i}{dt} = f(V_i, m_i, h_i, n_i, C_{m_i}, g_{K_i}...) \qquad (33)$$

This is now easily parallelizable using a vector computation of a form:

$$\frac{d\vec{V}}{dt} = f(\vec{V}, \vec{m}, \vec{h}, \vec{n}, \vec{C_m}, \vec{g_K}...) \qquad (34)$$

Thus we can do the calculations as:

$$\frac{d\mathbf{X'}}{dt} = \left[\frac{d\vec{V}}{dt}, \frac{d\vec{m}}{dt}, \frac{d\vec{h}}{dt}, \frac{d\vec{n}}{dt}\right] \qquad (35)$$

Implementation

Notice that the functions calculating gating dynamics and channel currents already are capable of vector input and output, so we do not need to change them. What we do need to change is the parameters which should now be vectors, the differential evaluation function which should now be designed to handle grouped parameters, and finally the initial condition.

```
n_n = 20 # number of simultaneous neurons to simulate
# parameters will now become n_n-vectors
C_m = [1.0]*n_n
g_K = [10.0]*n_n
E_K = [-95.0]*n_n
g_Na = [100]*n_n
E_Na = [50]*n_n
g_L = [0.15]*n_n
E_L = [-55.0]*n_n
```

```
10
11  # The state vector definition will change
12  def dXdt(X, t):
13      V = X[:1*n_n]        # First n_n values are Membrane Voltage
14      m = X[1*n_n:2*n_n]   # Next n_n values are Sodium Activation Gating
15      h = X[2*n_n:3*n_n]   # Next n_n values are Sodium Inactivation Gating
16      n = X[3*n_n:]        # Last n_n values are Potassium Gating
17      dVdt = (np.linspace(0,10,n_n)-I_Na(V, m, h)-I_K(V, n)-I_L(V))/ C_m
18      # Input current is linearly varied between 0 and 10
19      m0,tm,h0,th = Na_prop(V)
20      n0,tn = K_prop(V)
21      dmdt = - (1.0/tm)*(m-m0)
22      dhdt = - (1.0/th)*(h-h0)
23      dndt = - (1.0/tn)*(n-n0)
24      out = tf.concat([dVdt,dmdt,dhdt,dndt],0)
25      return out
26  y0 = tf.constant([-71]*n_n+[0,0,0]*n_n, dtype=tf.float64)
```

The output can be seen in Fig.   319

Quantifying the Firing Rates against Input Current   320

One way to quantify the firing rate is to perform a fourier analysis and find peak   321
frequency, but an easier way to find the rate is to see how many times it crosses a   322
threshold say 0 mV in a given time, here it is for 200ms = 0.2s, and find the rate. The   323
output can be seen in Fig.   324

```
1  rate = np.bitwise_and(state[:-1,:20]<0,state[1:,:20]>0).sum(axis=0)/0.2
```

Day 4: Neurons and Networks   325

In this section we simulate a network of neurons interacting via synapses. Each synapse   326
is defined by its own set of state variables and differential equations governing their   327
temporal evolution. There are different kinds of synapses - electrical and chemical   328
synapses. Electrical synapses are essentially physical conduits that allow the flow of ions 329
across connected neurons. Chemical synapses are more common in the brain and are   330
more complex than electrical synapses. When an action potential arrives at the axon   331
terminal, it leads to the opening of voltage-gated calcium channels. The incoming   332
calcium triggers neurotransmitter filled vescles to fuse with the axon terminal   333
membrane and release their cargo into the synaptic cleft. The neurotransmitters diffuse   334
across the cleft and open (or close) ion channels on the post-synaptic neuron. This can   335
cause a depolarization (increase in potential across the post-synaptic neuron's   336
membrane) that makes it easier for the neuron to spike or it can inhibit the neuron and   337
have the opposite effect. In some cases these effects are fast and direct —a   338
neurotransmitter binds to a receptor in the post-synaptic site that causes an influx or   339
eflux of ions and leads to a change in the membrane potential. The effect of a synapse   340
can also be indirect such that neurotransmitters invoke a second messenger cascade that   341
eventually leads to the opening or closing of ion channels in the post-synaptic neuron.   342
Here we model fast excitatory and inhibitory chemical synapses.The network of   343
interactions between neurons will be described by a connectivity matrix. Different   344
connectivity matrices describe the interactions due to different types of synapses.   345

Modelling Synapses

The synaptic current (I_{syn}) depends on the difference between its reversal potential (E_{syn}) and the actual value of the membrane potential (u). The synaptic current due to neurotransmitter release into the synaptic cleft following an action potential is given by,

$$I_{syn}(t) = g_{syn}(t)(u(t) - E_{syn}) \tag{36}$$

When the transmitter binds to postsynaptic receptors it causes a transient change in the conductance g_syn. To capture the dynamics of g_syn, one models the system using a simple kinetic model where the receptors can be either in an open or a closed state. The transition between the states is proportional to the concentration of the neurotransmitter $[T]$ in the cleft.

$$\text{C} \underset{\beta}{\overset{\alpha[T]}{\rightleftharpoons}} \text{O} \tag{37}$$

This may be rewritten in the form of a differential equation.

$$\frac{d[O]}{dt} = \alpha[T](1 - [O]) - \beta[O] \tag{38}$$

We can now describe the synaptic conductance $g_{syn}(t) = g_{max}[O](t)$, in terms of the maximal conductance g_{max} and a gating variable $[O]$, where $[O](t)$ is the fraction of open synaptic channels. α is the binding constant, β the unbinding constant and $(1 - [O])$ the fraction of closed channels where the neurotransmitter can bind. The functional form of T depends on the type and nature of the synapse. For cholinergic excitatory synapses, $[T]$ is given by,

$$[T]_{ach} = A\ \Theta(t_{max} + t_{fire} + t_{delay} - t)\ \Theta(t - t_{fire} - t_{delay}) \tag{39}$$

where, $\Theta(x)$ is the Heaviside step function, t_{fire} is the time of the last presynaptic spike, t_{delay} is the time delay from the time of the last spike to its effect on the postsynaptic neuron and t_{max} is the duration after the spike during which the transmitter remains in the synaptic cleft. For Fast GABAergic inhibitory synapses, we used the following equation,

$$[T]_{gaba} = \frac{1}{1 + e^{-\frac{V(t - t_{fire} - t_{delay}) - V_0}{\sigma}}} \tag{40}$$

Note that in order to solve the equation 38, we need to determine the time when the presynaptic neuron fired (t_{fire}). To account for these synaptic interactions between neurons we need to modify the RK4 integrator developed to simulate multiple independent Hodgkin-Huxley neurons.

Redesigning the Generalized TensorFlow Integrator

In this section we modify the integrator that we coded on Day 2 to account for interactions between neurons. This will require an additional variable that stores the time elapsed from the last presynaptic spike to calculate the equations 39, 40. In the modified code we will use the TensorFlow function `tf.where()` to efficiently assign the indices of neurons that have spiked and those that have not at each time point. To understand the usage and function of `tf.where()`, consider the following example. Say, you have a array `x` of 10 random numbers between 0 and 1. You want the output of the code to be another array of the same size as `x` such that the elements of the array are either -10 or 10 depending on whether the corresponding element in `x` is less or greater than 0.5. The function `tf.where(cond,a,b)` outputs an array with elements from `a` if the condition `cond` is `True` and from `b` if `cond` is `False`. See the example code below.

```
1   # create the Tensor with the random variables
2   x = tf.constant(np.random.uniform(size = (10,)),dtype=tf.float64)
3   # a list of 10s to select from if true
4   if_true = tf.constant(10*np.ones((10,)),dtype=tf.float64)
5   # a list of -10s to select from if false
6   if_false = tf.constant(-10*np.ones((10,)),dtype=tf.float64)
7   # perform the conditional masking
8   selection = tf.where(tf.greater(x,0.5),if_true,if_false)
9   with tf.Session() as sess:
10      x_out = sess.run(x)
11      selection_out = sess.run(selection)
12  # If x_out = [0.13 0.08 0.58 0.17 0.34 0.58 0.97 0.66 0.30 0.29 ],
13  # selection_out = [-10. -10.  10. -10. -10.  10.  10.  10. -10. -10.]
```

In order to determine whether a particular neuron fired, we introduce a new variable `fire_t` that stores the time of the last spike for each neuron.

We modify the code as follows:

1. The Integrator class that we defined earlier now requires two more properties, namely, the number of neurons (`n`) and firing threshold (`F_b`) of each of these neurons. We provide these inputs as arguments to the Integrator class.

2. The state vector will now have an additional `n` variables representing the firing times. These will not be updated by the step function (`_step_func`).

3. Inside the Integrator class, we have access to the values of the state variable and the change in the state variable since the last iteration. We use this to check if the voltages have crossed the firing threshold. The convention followed in this code is, the first `n` elements of the state vector are the membrane voltages while the last `n` elements are the time from the last spike for each of the neurons.

4. The differential update function ie. step_func takes except the last `n` values of the state variable and updates according to the differential equations specified in `func`. The last `n` variables are updated separately in `scan_func`. It checks if any neuron has crossed its firing threshold and updates the variable `fire_t` of the appropriate neurons with the current time.

The modifications to the RK4 code implemented earlier is shown below,

```
1   def integrate(self, func, y0, t):
2           time_delta_grid = t[1:] - t[:-1]
3           def scan_func(y, t_dt):
4               # recall the necessary variables
5               n_ = self.n_
6               F_b = self.F_b
7               t, dt = t_dt
8               # Differential updation
9               dy = self._step_func(func,t,dt,y) # Make code more modular.
10              dy = tf.cast(dy, dtype=y.dtype) # Failsafe
11              out = y + dy # the result after differential updation
12              # Use specialized Integrator vs Normal Integrator (n=0)
13              if n_>0:
14                  # Extract the last n variables for fire times
15                  fire_t = y[-n_:]
16                  # Change in fire_t if neuron didnt fire = 0
```

```
17                      l = tf.zeros(tf.shape(fire_t),dtype=fire_t.dtype)
18                      # Change in fire_t if neuron fired = Current-Last Fire
19                      l_ = t-fire_t
20                      # Check if previous Voltage is less than Threshold
21                      z = tf.less(y[:n_],F_b)
22                      # Check if Voltage is more than Threshold after update
23                      z_ = tf.greater_equal(out[:n_],F_b)
24                      df = tf.where(tf.logical_and(z,z_),l_,l)
25                      fire_t_ = fire_t+df # Update firing time
26                      return tf.concat([out[:-n_],fire_t_],0)
27                  else:
28                      return out
29          y = tf.scan(scan_func, (t[:-1], time_delta_grid),y0)
30          return tf.concat([[y0], y], axis=0)
31
32  def odeint(func, y0, t, n_, F_b):
33      t = tf.convert_to_tensor(t, preferred_dtype=tf.float64, name='t')
34      y0 = tf.convert_to_tensor(y0, name='y0')
35      tf_check_type(y0,t)
36      return _Tf_Integrator(n_, F_b).integrate(func,y0,t)
```

Implementing a network of Hodgkin-Huxley neurons

Recall, each Hodgkin Huxley Neuron in a network with n neurons has 4 dynamical variables V, m, n, h. Each of these variables were represented as n—dimensional vectors. Now we need to add some more state variables representing each synapse. The neuron receives excitatory and inhibitory inputs that are introduced as additional synaptic currents I_{ach} and I_{GABA}. Equation 25 now reads,

$$C_m \frac{dV}{dt} = I_{injected} - I_{Na} - I_K - I_L - I_{ach} - I_{gaba} \tag{41}$$

For each synapse, we will have Eq (38), Eq (39) and:

$$\frac{d[O]_{ach/gaba}}{dt} = \alpha(1 - [O]_{ach/gaba})[T]_{ach/gaba} - \beta[O]_{ach/gaba} \tag{42}$$

$$I_{ach/gaba}(t) = g_{max}[O]_{ach/gaba}(V - E_{ach/gaba}) \tag{43}$$

Synaptic Memory Management

As discussed earlier, there are atmost n^2 synapses of each type but at a time, unless the network is fully connected/very dense, mostly we need a very small subset of these synapses. We could, in principle, calculate the dynamics of all n^2 but it would be pointless. So have to devise a matrix based sparse-dense coding system for evaluating the dynamics of these variables and also using their values. This will reduce memory usage and minimize number of calculations at the cost of time for encoding and decoding into dense data from sparse data and vice versa. This is why we use a matrix approach, so that tensorflow can speed up the process.

Defining the Connectivity

Lets take a very simple 3 neuron network to test how the two types of synapses work. Let X_1 be an Excitatory Neuron and it forms Acetylcholine Synapses, X_2 be a Inhibitory Neuron and it forms GABAa Synapses and X_3 be an output neuron that doesn't synapse onto any neurons. Take the network of the form: $X_1 \rightarrow X_2 \rightarrow X_3$.

We create the connectivity matrices for both types of synapses. We need to define a convention for the ordering of connectivity. We set the presynaptic neurons as the column number, and the postsynaptic neurons as the row number.

Let X_1, X_2, X_3 be indexed by 0, 1 and 2 respectively. The Acetylcholine connectivity matrix takes the form of:

$$Ach_{n \times n} = \begin{bmatrix} 0 & 0 & 0 \\ 1 & 0 & 0 \\ 0 & 0 & 0 \end{bmatrix} \tag{44}$$

Similarly, the GABAa connectivity matrix becomes:

$$GABA_{n \times n} = \begin{bmatrix} 0 & 0 & 0 \\ 0 & 0 & 0 \\ 0 & 1 & 0 \end{bmatrix} \tag{45}$$

```
1   n_n = 3 # number of simultaneous neurons to simulate
2   # Acetylcholine
3   ach_mat = np.zeros((n_n,n_n))          # Ach Synapse Connectivity Matrix
4   ach_mat[1,0]=1
5   ## Parameters for Acetylcholine synapses ##
6   n_ach = int(np.sum(ach_mat))      # Number of Acetylcholine (Ach) Synapses
7   alp_ach = [10.0]*n_ach            # Alpha for Ach Synapse
8   bet_ach = [0.2]*n_ach             # Beta for Ach Synapse
9   t_max = 0.3                            # Maximum Time for Synapse
10  t_delay = 0                            # Axonal Transmission Delay
11  A = [0.5]*n_n                          # Synaptic Response Strength
12  g_ach = [0.35]*n_n                     # Ach Conductance
13  E_ach = [0.0]*n_n                      # Ach Potential
14  # GABAa
15  gaba_mat = np.zeros((n_n,n_n))         # GABAa Synapse Connectivity Matrix
16  gaba_mat[2,1] = 1
17  ## Parameters for GABAa synapses ##
18  n_gaba = int(np.sum(gaba_mat)) # Number of GABAa Synapses
19  alp_gaba = [10.0]*n_gaba        # Alpha for GABAa Synapse
20  bet_gaba = [0.16]*n_gaba        # Beta for GABAa Synapse
21  V0 = [-20.0]*n_n                       # Decay Potential
22  sigma = [1.5]*n_n                      # Decay Time Constant
23  g_gaba = [0.8]*n_n                     # fGABA Conductance
24  E_gaba = [-70.0]*n_n                   # fGABA Potential
25  ## Storing Firing Thresholds ##
26  F_b = [0.0]*n_n                        # Fire threshold
27  ## Store our input to each neuron as a n x timesteps matrix
28  ## called current_input, and extract value at each timepoint
29  def I_inj_t(t):
30      # Turn indices to integer and extract from matrix
31      index = tf.cast(t/epsilon,tf.int32)
32      return tf.constant(current_input.T,dtype=tf.float64)[index]
```

Working with Sparse-Dense Dynamics

For performing the dynamical updates for the synapses, we need only as many variables as the number of synapse x number of equations required for each synapse. Here our synapse models require only one dynamical variable ie. open fraction [O] which we will store as an k-vector where k is the number of synapses.

We will need to work with this [O] vector on two seperate instances and each time we will have to go to a dense form to speed up computation.

A. Calculation of Synaptic Currents

The formula for I_{syn} is given by:

$$I_{syn} = \sum_{presynaptic} g_{syn}[O](V - E_{syn}) \tag{46}$$

The best way to represent this calculation is to use the connectivity matrix $\mathbf{C}$ for the synapses and the open fraction vector $[\vec{O}]$ to create an open fraction matrix $\mathbf{O}$ and perform the following computations.

$$\mathbf{C} = \begin{bmatrix} 0 & 1 & ... & 0 \\ 0 & 0 & ... & 1 \\ ... & ... & ... & 1 \\ 1 & 0 & 0 & 0 \end{bmatrix} \tag{47}$$

$$[\vec{O}] = [O_1, O_2...O_k] \tag{48}$$

$$\mathbf{O} = \begin{bmatrix} 0 & O_1 & ... & 0 \\ 0 & 0 & ... & O_a \\ ... & ... & ... & O_b \\ O_k & 0 & 0 & 0 \end{bmatrix} \tag{49}$$

$$[\vec{I_{syn}}] = \sum_{columns} \mathbf{O} \diamond (\vec{g}_{syn} \odot (\vec{V} - \vec{E}_{syn})) \tag{50}$$

where $\diamond$ is columnwise multiplication and $\odot$ is elementwise multiplication. $[\vec{I_{syn}}]$ is now the total synaptic current input to the each of the neurons.

Algorithm for Synaptic Currents

1. Firstly we need to convert from the sparse $[O]$ vector to the dense $\mathbf{O}$ matrix. TensorFlow does not allow to make changes to a defined tensor directly, thus we create a n^2 vector TensorFlow variable o_ which we will later reshape to a $n \times n$ matrix.

2. We then flatten the synaptic connectivity matrix and find the indices where there is a connection. For this we use the boolean mask function to choose the correct k (total number of synapses) indices from the range 1 to n^2 and store in the variable ind.

3. Using the scatter_update function of TensorFlow, we fill the correct indices of the variable o_ that we created with the values of open fraction from the [O] vector.

4. We now reshape the vector as a $n \times n$ matrix. Since python stores matrices as array of arrays, with each row as an inner array, for performing columnswise multiplication, the easiest way is to tranpose the matrix, so that each column is the inner array, perform element wise multiplication with each inner array and apply transpose again.

5. Finally using reduce_sum, we sum over the columns to get our I_{syn} vector.

```
## Acetylcholine Synaptic Current ##
def I_ach(o,V):
    o_ = tf.Variable([0.0]*n_n**2,dtype=tf.float64)
    ind = tf.boolean_mask(tf.range(n_n**2),ach_mat.reshape(-1) == 1)
    o_ = tf.scatter_update(o_,ind,o)
    o_ = tf.transpose(tf.reshape(o_,(n_n,n_n)))
    return tf.reduce_sum(tf.transpose((o_*(V-E_ach))*g_ach),1)
## GABAa Synaptic Current ##
def I_gaba(o,V):
    o_ = tf.Variable([0.0]*n_n**2,dtype=tf.float64)
    ind = tf.boolean_mask(tf.range(n_n**2),gaba_mat.reshape(-1) == 1)
    o_ = tf.scatter_update(o_,ind,o)
    o_ = tf.transpose(tf.reshape(o_,(n_n,n_n)))
    return tf.reduce_sum(tf.transpose((o_*(V-E_gaba))*g_gaba),1)
## Other Currents remain the same ##
```

B. Updation of Synaptic Variables

For the updation, the first we need to calculate the values of the presynaptic activation [T] for both types of synapses. We will essentially calculate this for each neuron and then redirect the values to the correct post synaptic neuron. Recall:

$$[T]_{ach} = A \; \Theta(t_{max} + t_{fire} + t_{delay} - t) \; \Theta(t - t_{fire} - t_{delay}) \tag{51}$$

$$[T]_{gaba} = \frac{1}{1 + e^{-\frac{V(t) - V_0}{\sigma}}} \tag{52}$$

Thus $[T]_{ach}$ is function of the last firing time, and $[T]_{gaba}$ depends on the presynaptic voltage. Once we calculate the values of [T]-vector for both types of synapse, we need to redirect them to the correct synapses in a sparse $k \times 1$ vector form.

Algorithm for Dynamics

1. For $[T]_{ach}$, use a boolean logical_and function to check is the current timepoint t is greater than the last fire time (fire_t) + delay (t_delay) and less than last fire time (fire_t) + delay (t_delay) + activation length (t_max) for each neuron as a vector. Use the result of these boolean operations to choose between zero or an constant A. This serves as the heaviside step function. For $[T]_{gaba}$, simply use the V vector to determine T.

2. For making the sparse vector, we follow as two step process. First we multiply each row of the connectivity matrices $\mathbf{C}$ with the respective [T] vector to get a

activation matrix $\mathbf{T}$, and then we just flatten $\mathbf{T}$ and $\mathbf{C}$ and, using
tf.boolean_mask, remove all the zeros from $\mathbf{T}$ to get a $k \times 1$ vector which now
stores the presynaptic activation for each of the synapses where $k = n_{gaba}$ or n_{ach}.

3. Calculate the differential change in the open fractions (OF) using the $k \times 1$ vector.

```python
def dXdt(X, t):
    V = X[:1*n_n]          # First n_n: Membrane Voltage
    m = X[1*n_n:2*n_n]     # Next n_n: Sodium Activation Gating
    h = X[2*n_n:3*n_n]     # Next n_n: Sodium Inactivation Gating
    n = X[3*n_n:4*n_n]     # Next n_n: Potassium Gating
    # Next n_ach and n_gaba: Ach and GABAa Open Fraction respectively
    o_ach = X[4*n_n : 4*n_n + n_ach]
    o_gaba = X[4*n_n + n_ach : 4*n_n + n_ach + n_gaba]
    fire_t = X[-n_n:]     # Last n_n: last fire times
    dVdt = (I_inj_t(t)-I_Na(V, m, h)-I_K(V, n)-
                    I_L(V)-I_ach(o_ach,V)-I_gaba(o_gaba,V))/C_m
    ## Updation for gating variables ##
    m0,tm,h0,th = Na_prop(V)
    n0,tn = K_prop(V)
    dmdt = - (1.0/tm)*(m-m0)
    dhdt = - (1.0/th)*(h-h0)
    dndt = - (1.0/tn)*(n-n0)
    ## Updation for o_ach ##
    A_ = tf.constant(A,dtype=tf.float64)
    Z_ = tf.zeros(tf.shape(A_),dtype=tf.float64)
    T_ach = tf.where(tf.logical_and(tf.greater(t,fire_t+t_delay),
                            tf.less(t,fire_t+t_max+t_delay)),A_,Z_)
    T_ach = tf.multiply(tf.constant(ach_mat,dtype=tf.float64),T_ach)
    T_ach = tf.boolean_mask(tf.reshape(T_ach,(-1,)),
                            ach_mat.reshape(-1) == 1)
    do_achdt = alp_ach*(1.0-o_ach)*T_ach - bet_ach*o_ach
    ## Updation for o_gaba ##
    T_gaba = 1.0/(1.0+tf.exp(-(V-V0)/sigma))
    T_gaba = tf.multiply(tf.constant(gaba_mat,dtype=tf.float64),T_gaba)
    T_gaba = tf.boolean_mask(tf.reshape(T_gaba,(-1,)),
                            gaba_mat.reshape(-1) == 1)
    do_gabadt = alp_gaba*(1.0-o_gaba)*T_gaba - bet_gaba*o_gaba
    ## Updation for fire times ##
    dfdt = tf.zeros(tf.shape(fire_t),dtype=fire_t.dtype) # no change
    out = tf.concat([dVdt,dmdt,dhdt,dndt,do_achdt,do_gabadt,dfdt],0)
    return out
```

Defining the Gating Variable Updation Function and the Initial Conditions

Like earlier, we again define the function that returns us the values of τ_m, τ_h, τ_n, m_0, h_0, n_0 and we prepare the parameters and initial conditions.

Note: If we initialize the last firing time as 0, then the second neuron X_2 will get an EPSP immediately after the start of the simulation. To avoid this the last firing time should be initialized to a large negative number ¿= the length of the simulation.

```python
# The initialization of the Parameters and Gating Variable
# Updation Function remain the same.
```

```
 3
 4   # Initialize the State Vector
 5   y0 = tf.constant([-71]*n_n+[0,0,0]*n_n+[0]*n_ach+[0]*n_gaba+
 6                              [-9999999]*n_n,dtype=tf.float64)
```

Creating the Current Input and Run the Simulation

We will run an 700 ms simulation with 100ms current injection at neuron X_1 of
increasing amplitude with 100ms gaps in between.

```
 1   current_input= np.zeros((n_n,t.shape[0]))
 2   current_input[0,int(100/epsilon):int(200/epsilon)] = 2.5
 3   current_input[0,int(300/epsilon):int(400/epsilon)] = 5.0
 4   current_input[0,int(500/epsilon):int(600/epsilon)] = 7.5
 5   state = odeint(dXdt,y0,t,n_n,F_b)
 6   with tf.Session() as sess:
 7       # Since we are using variables we have to initialize them
 8       tf.global_variables_initializer().run()
 9       state = sess.run(state)
```

The output is plotted in Fig.
We can see that the current injection triggers the firing of action potentials with
increasing frequency with increasing current. Also we see that as soon as the first
neuron X_1 crosses its fireing threshold, an EPSP is triggered in the next neuron X_2
causing a firing with a slight delay from the firing of X_1. Finally, as the second neuron
depolarizes, we see a corresponding hyperpolarization in the next neuron X_3 caused by
an IPSP. We can also plot the dynamics of the channels itself by plotting o_ach and
o_gaba which are the 5th and 4th last elements respectively (Fig). Thus we are now
capable of making complex networks of neurons with both excitatory and inhibitory
connections.

Day 5: Optimal Mind Control

Now that we can simulate a model of a network of conductance-based neurons, we
discuss the limitations of our algorithm and try to find solutions to the problems.

Memory Management

This TensorFlow implementation allows us to make our simulation code not only easier
to read but also makes it highly parallizable and scalable across a variety of
computational devices. But there are some major limitations to our system. The
biggest of these issues is that despite making the simulation faster, this implementation
makes it very memory intensive.
The iterators in TensorFlow do not follow the normal process of Memory Allocation
and Garbage Collection. Since, TensorFlow is designed to work on sophisticated
hardware like GPUs, TPUs and distributed platforms, the memory allocation is done
during TensorFlow session adaptively, but the memory is NOT cleared until the python
kernel has stopped execution.
The memory used increases linearly with time as the state matrix is computed
recursively by the Scan function. The maximum memory used by the computational
graph is 2 times the total state matrix size at the point when the computation finishes
and copies the final data into the memory. Larger and more complicated the network
and longer the simulation, larger the matrix and then, each run is limited by the total

available memory. As we increase the number of neurons (n) and scale the degree ₅₂₀ distributions, the size of the memory used grows $O(n^2)$ as the number of differential ₅₂₁ equations grows as a square of the number of neurons. Similarly, as we increase the ₅₂₂ length of the simulation(L), the memory used grows $O(L)$. Which is why for a system ₅₂₃ with a limited memory of K bytes, The length of a given simulation (L timesteps) of a ₅₂₄ given network (N differential equations) with 64-bit floating-point precision will follow: ₅₂₅

$$2 \times 64 \times L \times N = K \tag{53}$$

That is, for any given network, our maximum simulation length is limited. One way ₅₂₆ to improve our maximum length is to divide the simulation into smaller batches. There ₅₂₇ will be a small queuing time between batches, which will slow down our code by a small ₅₂₈ amount but we will be able to simulate longer times. Thus, if we split the simulation ₅₂₉ into K sequential batches, the maximum memory for the simulation becomes $(1 + \frac{1}{K})$ ₅₃₀ times the total matrix size. Thus the memory relation becomes: ₅₃₁

$$\left(1 + \frac{1}{K}\right) \times 64 \times L \times N = K \tag{54}$$

This way, we can maximize the length of out simulation that we can run in a single ₅₃₂ python kernel. ₅₃₃

Let us implement this batch system for our 3 neuron feed-forward model. ₅₃₄

Implementing the Model ₅₃₅

To improve out readability and make our code more modular we seperate the integrator ₅₃₆ into a independent import module. Take the integrator code that we developed in the ₅₃₇ last day and place it in a file called "tf_integrator.py". Make sure the file is present in ₅₃₈ the same directory as the implementation of the model. ₅₃₉

Note: If you are using Jupyter Notebook, remember to remove the %matplotlib ₅₄₀ inline command as it is specific to jupyter. ₅₄₁

Importing tf_integrator and other requirements ₅₄₂

Once the Integrator is saved in tf_integrator.py in the same directory as the Notebook, ₅₄₃ we can start importing the essentials including the integrator. ₅₄₄

```
1  import tensorflow as tf
2  import numpy as np
3  import tf_integrator as tf_int
4  import matplotlib.pyplot as plt
5  import seaborn as sns
```

For implementing a Batch system, we do not need to change how we construct our ₅₄₅ model only how we execute it. ₅₄₆

Splitting Time Series into independent batches and Run Each Batch ₅₄₇ Sequentially ₅₄₈

Since we will be dividing the computation into batches, we have to split the time array ₅₄₉ into batches which will be passed to the each successive call to the integrator. For each ₅₅₀ new call, the last state vector of the last call will be the new initial condition. ₅₅₁

In np.array_split(), the split edges are present in only one array and since our initial ₅₅₂ vector to successive calls is corresponding to the last output our first element in the ₅₅₃

later time array should be the last element of the previous output series. Thus, we append the last time to the beginning of the current time array batch. 554 555

```python
# Define the Number of Batches
n_batch = 2
# Split t array into batches using numpy
t_batch = np.array_split(t,n_batch)
# Iterate over the batches of time array
for n,i in enumerate(t_batch):
    # Inform start of Batch Computation
    print("Batch",(n+1),"Running...",end="")
    # Re-adjusting edges
    if n>0:
        i = np.append(i[0]-sim_res,i)
    # Set state_vector as the initial condition
    init_state = tf.constant(state_vector, dtype=tf.float64)
    # Create the Integrator computation graph
    tensor_state = tf_int.odeint(dXdt, init_state, i, n_n, F_b)
    # Initialize variables and run session
    with tf.Session() as sess:
        tf.global_variables_initializer().run()
        state = sess.run(tensor_state)
        sess.close()
    # Reset state_vector as the last element of output
    state_vector = state[-1,:]
    # Save the output of the simulation to a binary file
    np.save("part_"+str(n+1),state)
    # Clear output
    state=None
    print("Finished")
```

Putting the Output Together 556

The output from our batch implementation is a set of binary files that store subparts of our total simulation. To get the overall output we have to stitch them back together. 557 558

```python
overall_state = []
# Iterate over the generated output files
for n,i in enumerate(["part_"+str(n+1)+".npy" for n in range(n_batch)]):
    # Since the first element in the series was the last output,
    # we remove them
    if n>0:
        overall_state.append(np.load(i)[1:,:])
    else:
        overall_state.append(np.load(i))
# Concatenate all the matrix to get a single state matrix
overall_state = np.concatenate(overall_state)
```

By this method, we have maximized the usage of our available memory but we can go further and develop a method to allow indefinitely long simulation. The issue behind this entire algorithm is that the memory is not cleared until the python kernel finishes. One way to overcome this is to save the parameters of the model (such as connectivity matrix) and the state vector in a file, and start a new python kernel from a python script to compute successive batches. This way after each large batch, the memory gets 559 560 561 562 563 564

cleaned. By combining the previous batch implementation and this system, we can   565
maximize our computability.   566

Implementing a Runner and a Caller   567

Firstly, we have to create an implementation of the model that takes in previous input   568
as current parameters. Thus, we create a file, which we call "run.py" that takes an   569
argument ie. the current batch number. The implementation for "run.py" is mostly   570
same as the above model but there is a small difference.   571

When the batch number is 0, we initialize all variable parameters and save them, but   572
otherwise we use the saved values. The parameters we save include: Acetylcholine   573
Matrix, GABAa Matrix and Final/Initial State Vector. It will also save the files with   574
both batch number and sub-batch number listed.   575

The time series will be created and split initially by the caller, which we call   576
"call.py", and stored in a file. Each execution of the Runner will extract its relevant   577
time series and compute on it.   578

Implementing the Caller code   579

The caller will create the time series, split it and use python subprocess module to call   580
"run.py" with appropriate arguments. The code for "call.py" is given below.   581

```
1   from subprocess import call
2   import numpy as np
3   total_time = 1000
4   n_splits = 2
5   time = np.split(np.arange(0,total_time,0.01),n_splits)
6   # Append the last time point to the beginning of the next batch
7   for n,i in enumerate(time):
8       if n>0:
9           time[n] = np.append(i[0]-0.01,i)
10  np.save("time",time)
11  # call successive batches with a new python subprocess
12  # and pass the batch number
13  for i in range(n_splits):
14      call(['python','run.py',str(i)])
15  print("Simulation Completed.")
```

Implementing the Runner code   582

"run.py" is essentially identical to the batch-implemented model we developed earlier   583
with the changes described below: 584

```
1   # Additional Imports #
2   import sys
3   # Duration of Simulation #
4   # Replace t = np.arange(0,sim_time,sim_res) by
5   t = np.load("time.npy")[int(sys.argv[1])] # get first argument to run.py
6   # Connectivity Matrix Definitions #
7   if sys.argv[1] == '0':
8       ach_mat = np.zeros((n_n,n_n)) # Ach Synapse Connectivity Matrix
9       ach_mat[1,0]=1
10      # If connectivity is random, once initialized it will be the same.
11      np.save("ach_mat",ach_mat)
```

```
12  else:
13      ach_mat = np.load("ach_mat.npy")
14  if sys.argv[1] == '0':
15      gaba_mat = np.zeros((n_n,n_n)) # GABAa Synapse Connectivity Matrix
16      gaba_mat[2,1] = 1
17      # If connectivity is random, once initialized it will be the same.
18      np.save("gaba_mat",gaba_mat)
19  else:
20      gaba_mat = np.load("gaba_mat.npy")
21  # Current Input Definition #
22  if sys.argv[1] == '0':
23      current_input= np.zeros((n_n,int(sim_time/sim_res)))
24      current_input[0,int(100/sim_res):int(200/sim_res)] = 2.5
25      current_input[0,int(300/sim_res):int(400/sim_res)] = 5.0
26      current_input[0,int(500/sim_res):int(600/sim_res)] = 7.5
27      np.save("current_input",current_input)
28  else:
29      current_input = np.load("current_input.npy")
30  # State Vector Definition #
31  if sys.argv[1] == '0':
32      state_vector = [-71]*n_n+[0,0,0]*n_n+[0]*n_ach+[0]*n_gaba
33                                          +[-9999999]*n_n
34      state_vector = np.array(state_vector)
35      state_vector = state_vector + 0.01*state_vector
36                          *np.random.normal(size=state_vector.shape)
37      np.save("state_vector",state_vector)
38  else:
39      state_vector = np.load("state_vector.npy")
40  # Saving of Output #
41  # Replace np.save("part_"+str(n+1),state) by
42  np.save("batch"+str(int(sys.argv[1])+1)+"_part_"+str(n+1),state)
```

Combining all Data 585

Just like we merged all the batches, we merge all the sub-batches and batches. 586

```
1   overall_state = []
2   # Iterate over the generated output files
3   for n,i in enumerate(["batch"+str(x+1) for x in range(n_splits)]):
4       for m,j in enumerate(["_part_"+str(x+1) for x in range(n_batch)]):
5           # Since the first element in the series was the last output,
6           # we remove them
7           if n>0 and m>0:
8               overall_state.append(np.load(i+j+".npy")[1:,:])
9           else:
10              overall_state.append(np.load(i+j+".npy"))
11  # Concatenate all the matrix to get a single state matrix
12  overall_state = np.concatenate(overall_state)
```

Other Limitations 587

With the runner-caller implementation, we have solved the issue of memory limitation. 588
But there are some issues we havent solved yet. 589

1. **Delay Dependent Activity:** If the implementation of an model requires a delay differential equation (DDEs) that accesses the value of the parameters at a earlier time, we cannot use this programming paradigm directly as it has only, at most, one timestep memory. This is also a limitation of using the scan function. Exploring alternatives to tf.scan() and other techniques of solving DDEs is an option.

2. **Optimal Distributed Computing:** To maximize the utility of distributed systems, the algorithm needs to be capable of splitting the computation into tasks that can be run parallely (simultaneously) in different computational units. Since we are working with numerical integration, the results of the last step need to be known to perform the next computation. This, in a way limits out ability to use distributed computing. TensorFlow automatically tries to utilize all available resources efficiently but it may not be the optimal usage. One way to enforce distribution is to manually distribute parts of the computational graph to specialized computing units. For example, the actual session can be forced to be executed on the server with maximum memory but the sub-computations such as the RK4 integration steps can be distributed to servers with more GPUs.

Discussion

Conclusion

Supporting information

S1 Fig. Bold the title sentence. Add descriptive text after the title of the item (optional).

Acknowledgments

References

1. Conant GC, Wolfe KH. Turning a hobby into a job: how duplicated genes find new functions. Nat Rev Genet. 2008 Dec;9(12):938–950.

2. Ohno S. Evolution by gene duplication. London: George Alien & Unwin Ltd. Berlin, Heidelberg and New York: Springer-Verlag.; 1970.

3. Magwire MM, Bayer F, Webster CL, Cao C, Jiggins FM. Successive increases in the resistance of Drosophila to viral infection through a transposon insertion followed by a Duplication. PLoS Genet. 2011 Oct;7(10):e1002337.